

Rohan Dubey

Hyderabad, India | +91-7093123447 | rohankumardubey497@gmail.com | github.com/rohankumardubey | rohankumardubey.github.io

SUMMARY

Senior Data Engineer (Java, Python, Go, Scala, C++) focused on real-time, low-latency distributed systems and cost-efficient data platforms on AWS. 7+ years building production pipelines with Flink, Kafka/MSK, Kinesis Data Analytics, Spark, and DynamoDB/S3. Proven impact on cost savings, performance tuning, and reliability for business-critical pipelines.

CORE STRENGTHS

- Languages: Java, Python, Go, Scala, C++, SQL, Shell.
- Streaming & Batch: Apache Flink, Apache Kafka (MSK), Spark, Kinesis Data Analytics (KDA), Airflow
- Storage & Databases: AWS DynamoDB, PostgreSQL, MySQL, Delta/Parquet/ORC, S3
- Infra & Cloud: AWS (S3, Glue, EMR, KDA, MSK, Redshift, Athena), Kubernetes, Docker
- Data Platforms & Tools: Dremio, Snowflake, NiFi/Registry, GitHub/GitLab, CI/CD
- Specialties: Streaming design, backpressure and checkpointing, schema evolution, serialization (Kryo/Avro/Protobuf), performance and cost optimization, observability

EXPERIENCE

PokerStars (Flutter Entertainment) - Senior Data Engineer, Streaming AI

Aug 2023 - Present, Hyderabad

- Designed real-time processing on AWS (Flink on Kinesis Data Analytics + MSK + DynamoDB + S3); cut hot-path IO and saved approx USD 600K per year via DynamoDB modeling, payload compaction, and S3 layout tuning.
 - Built exactly-once sinks with checkpoint alignment and idempotent upserts; tuned RocksDB state and JVM for reliable recovery and low tail latency.
-

Experian PLC - Senior Data Engineer

Mar 2023 - Aug 2023, Hyderabad

- Delivered regulated pipelines with secure ingestion, data quality gates, and lineage; accelerated analytics readiness for clients.
 - Standardized batch and streaming jobs with reproducible configs and monitoring.
-

Model N - Member of Technical Staff 3

Jun 2021 - Mar 2023, Hyderabad

- Built event-driven analytics with Kafka, Spark, and Delta Lake; exposed queries via Dremio and REST services.
 - Improved query performance using partitioning, Z-ordering, predicate pushdown, and compaction; reduced compute and storage costs.
-

Carelon Global Solutions - Software Engineer

Aug 2020 - Jun 2021, Hyderabad

- Integrated Medicare/Medicaid datasets and SQL analytics to improve revenue capture and reduce overhead.
-

Legato Health Technologies - Associate Software Engineer

Jul 2018 - Aug 2020, Hyderabad

- Delivered optimized ETL on Teradata and Informatica; standardized SLAs and validations.
-

SELECTED PROJECTS

- GoXStream: Flink-inspired stream processor in Go with operator graph, checkpoints, and Kafka/file/db connectors (github.com/rohankumardubey/goxstream).
- StreamNest: Minimal Kafka-like broker in Go with log segmentation, producer/consumer CLI, and replication primitives (github.com/rohankumardubey/StreamNest).

EDUCATION

Mahatma Gandhi Institute of Technology, Hyderabad - B.Tech (Electrical), GPA 4.0/4.0 (2014 - 2018)

KEYWORDS (ATS)

Distributed Systems, Real-time Streaming, Low Latency, Apache Kafka, Apache Flink, Apache Spark, AWS (S3, DynamoDB, MSK, Kinesis, Lambda), Kubernetes, Docker, CI/CD, Java, C++, Go, Python, SQL, Parquet/ORC, Observability, Exactly-Once, Checkpointing, Data Modeling, Cost Optimization, Data Lake, MapReduce, HBase, Cassandra, NiFi, Athena, PostgreSQL, AWS Glue, AWS Secret Manager, Git, Maven, Apache Iceberg, Airbyte, Jenkins, Fabric8, Bitbucket, Teradata, GCP, Big Query.